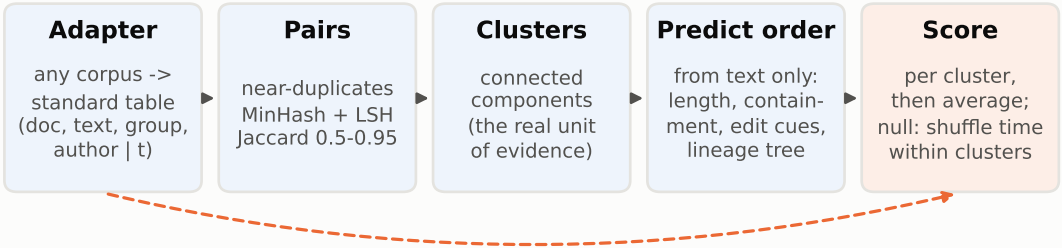
Adapter
any corpus -> standard table (doc, text, group, author | t)
Pairs
near-duplicates MinHash + LSH Jaccard 0.5-0.95
Clusters
connected components (the real unit of evidence)
Predict order
from text only: length, containment, edit cues, lineage tree
Score
per cluster, then average; null: shuffle time within clusters
true time t bypasses the method: only the scoring step (src/evaluate.py) reads it